# Homework 5

# Jacob Adams

# Table of contents

																					4
Question 1			 																		2
Question 2																					12
Question 3																					20

# Important

Please read the instructions carefully before submitting your assignment.

- 1. This assignment requires you to only upload a PDF file on Canvas
- 2. Don't collapse any code cells before submitting.
- 3. Remember to make sure all your code output is rendered properly before uploading your submission.

Please add your name to the author information in the frontmatter before submitting your assignment

In this assignment, we will explore decision trees, support vector machines and neural networks for classification and regression. The assignment is designed to test your ability to fit and analyze these models with different configurations and compare their performance.

We will need the following packages:

```
packages <- c(
  "tibble",
  "dplyr",
  "readr",</pre>
```

```
"tidyr",
  "purrr",
  "broom",
  "magrittr",
  "corrplot",
  "caret",
  "rpart",
  "rpart.plot",
  "e1071",
  "torch",
 "luz"
#renv::install(packages)
sapply(packages, require, character.only=T)
```

# Question 1



• 60 points

Prediction of Median House prices

#### 1.1 (2.5 points)

The data folder contains the housing.csv dataset which contains housing prices in California from the 1990 California census. The objective is to predict the median house price for California districts based on various features.

Read the data file as a tibble in R. Preprocess the data such that:

- 1. the variables are of the right data type, e.g., categorical variables are encoded as factors
- 2. all column names to lower case for consistency
- 3. Any observations with missing values are dropped

```
path <- "data/housing.csv"</pre>
df <- read_csv(path)</pre>
```

```
Rows: 20640 Columns: 10
-- Column specification -----
Delimiter: ","
chr (1): ocean_proximity
dbl (9): longitude, latitude, housing_median_age, total_rooms, total_bedroom...

i Use `spec()` to retrieve the full column specification for this data.
i Specify the column types or set `show_col_types = FALSE` to quiet this message.

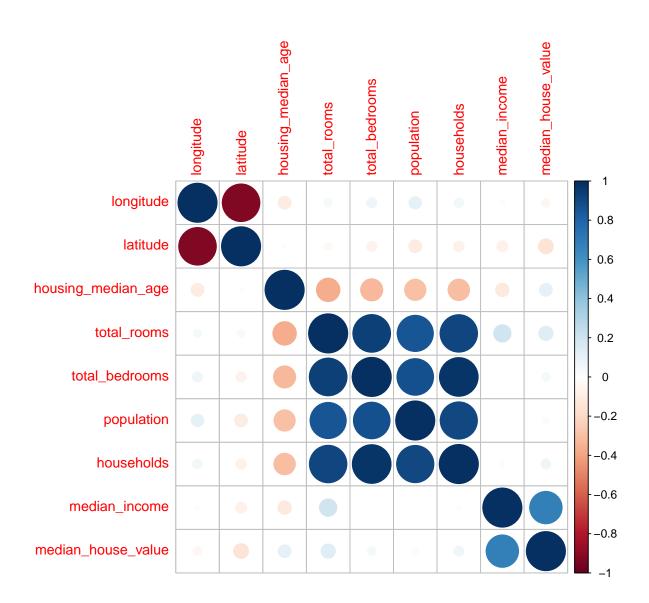
df <- df %>%
    mutate(across(where(is.character), as.factor)) %>%
    rename_all(tolower) %>%
    drop_na()
```

# 1.2 (2.5 points)

Visualize the correlation matrix of all numeric columns in df using corrplot()

```
numeric_data <- df %>% select_if(is.numeric)
correlation_matrix <- cor(numeric_data)

corrplot(correlation_matrix)</pre>
```



# 1.3 (5 points)

Split the data df into df\_train and df\_split using test\_ind in the code below:

```
set.seed(42)
test_ind <- sample(</pre>
```

```
1:nrow(df),
floor( nrow(df)/10 ),
replace=FALSE
)

df_test <- df[-test_ind,]
df_train <- df[test_ind,]</pre>
```

## 1.4 (5 points)

Fit a linear regression model to predict the median\_house\_value:

- latitude
- longitude
- housing\_median\_age
- total\_rooms
- total\_bedrooms
- population
- median\_income
- ocean\_proximity

Interpret the coefficients and summarize your results.

```
lm_fit <- lm(median_house_value ~ latitude + longitude + housing_median_age + total_rooms + rounds + round
```

```
latitude
                                      3.345e+03 -10.515 < 2e-16 ***
                          -3.517e+04
longitude
                          -3.746e+04
                                      3.376e+03 -11.096 < 2e-16 ***
                           1.144e+03
                                      1.426e+02
                                                   8.027 1.67e-15 ***
housing_median_age
total rooms
                                                 -3.484 0.000504 ***
                          -8.063e+00
                                      2.314e+00
total bedrooms
                           1.568e+02
                                      1.151e+01
                                                  13.622
                                                          < 2e-16 ***
population
                          -3.821e+01
                                      3.271e+00 -11.682
                                                          < 2e-16 ***
median income
                           3.981e+04
                                      1.028e+03
                                                  38.709
                                                          < 2e-16 ***
ocean_proximityINLAND
                          -3.001e+04
                                      5.671e+03
                                                  -5.292 1.34e-07 ***
ocean proximityISLAND
                                      6.837e+04
                                                   3.002 0.002712 **
                           2.053e+05
ocean_proximityNEAR BAY
                          -1.759e+04
                                      5.968e+03
                                                  -2.947 0.003243 **
                                                   0.217 0.828219
ocean_proximityNEAR OCEAN
                                      4.901e+03
                           1.064e+03
```

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 68200 on 2031 degrees of freedom Multiple R-squared: 0.6711, Adjusted R-squared: 0.6693 F-statistic: 376.7 on 11 and 2031 DF, p-value: < 2.2e-16

Interpretation of Coefficients: Latitude: When all covariates are held constant, there is a 2.539e+04 decrease in median house value for every unit increase in latitude. Longitude: When all covariates are held constant, there is a 2.681e+04 decrease in median house value for every unit increase in longitude. Median House Age: When all covariates are held constant, there is a 1.074e+03 increase in median house value for every unit increase in median house age. Total Rooms: When all covariates are held constant, there is a 6.159e+00 decrease in median house value for every unit increase in total rooms. Total Bedrooms: When all covariates are held constant, there is a 1.353e+02 increase in median house value for every unit increase in total bedrooms. Population: When all covariates are held constant, there is a 3.413e+01 decrease in median house value for every unit increase in population. Median Income: When all covariates are held constant, there is a 3.936e+04 increase in median house value for every unit increase in population. Ocean proximity: When all covariates are held constant, there is a -4.018e+04, 1.324e+05, -2.522e+03, and 4.349e+03 change in median house value when a resident lives inland, on an island, near bay, and near ocean respectively. It should be noted living near a bay is the only non-significant predictor compared to the rest of the covariates.

#### 1.5 (5 points)

Complete the rmse function for computing the Root Mean-Squared Error between the true y and the predicted yhat, and use it to compute the RMSE for the regression model on df\_test

```
rmse <- function(y, yhat) {
   sqrt(mean((y - yhat)^2))
}

lm_predictions <- predict(lm_fit, newdata = df_test)
rmse_vals <- rmse(df_test$median_house_value, lm_predictions)</pre>
```

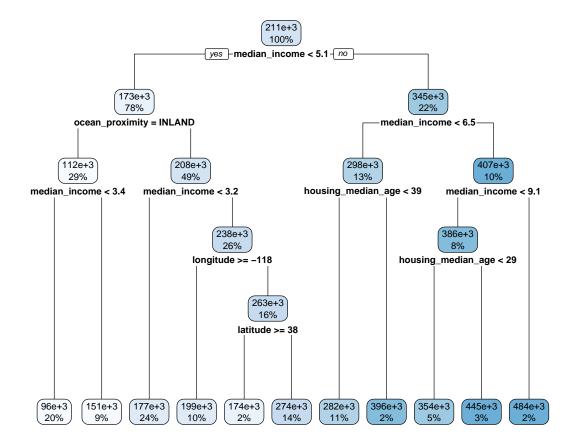
1.6 (5 points)

Fit a decision tree model to predict the median\_house\_value using the same predictors as in 1.4. Use the rpart() function.

```
rpart_fit <- rpart(median_house_value ~ latitude + longitude + housing_median_age + total_ro
rpart_predictions <- predict(rpart_fit, newdata = df_test)</pre>
```

Visualize the decision tree using the rpart.plot() function.

```
library(rpart.plot)
rpart.plot(rpart_fit, box.palette = "Blues")
```



Report the root mean squared error on the test set.

rpart\_predictions\_rmse <- rmse(df\_test\$median\_house\_value, rpart\_predictions)
print(rpart\_predictions\_rmse)</pre>

[1] 74346.43

### 1.7 (5 points)

Fit a support vector machine model to predict the median\_house\_value using the same predictors as in 1.4. Use the svm() function and use any kernel of your choice. Report the root mean squared error on the test set.

```
svm_fit <- svm(median_house_value ~ latitude + longitude + housing_median_age + total_rooms
svm_predictions <- predict(svm_fit, newdata = df_test)

svm_rmse <- rmse(svm_predictions, df_test$median_house_value)
svm_rmse

[1] 61064.47</pre>
```

# 1.8 (25 points)

Initialize a neural network model architecture:

```
NNet <- nn_module(</pre>
    initialize = function(p, q1, q2, q3){
      self$hidden1 <- nn_linear(p,q1)</pre>
      self$hidden2 <- nn_linear(q1, q2)</pre>
      self$hidden3 <- nn_linear(q2, q3)</pre>
      self$output <- nn_linear(q3, 1)</pre>
      self$activation <- nn_relu()</pre>
      self$sigmoid <- nn_sigmoid()</pre>
    },
    forward = function(x){
      x %>%
         self$hidden1() %>% self$activation() %>%
         self$hidden2() %>% self$activation() %>%
         self$hidden3() %>% self$activation() %>%
         self$output() %>% self$sigmoid()
    }
```

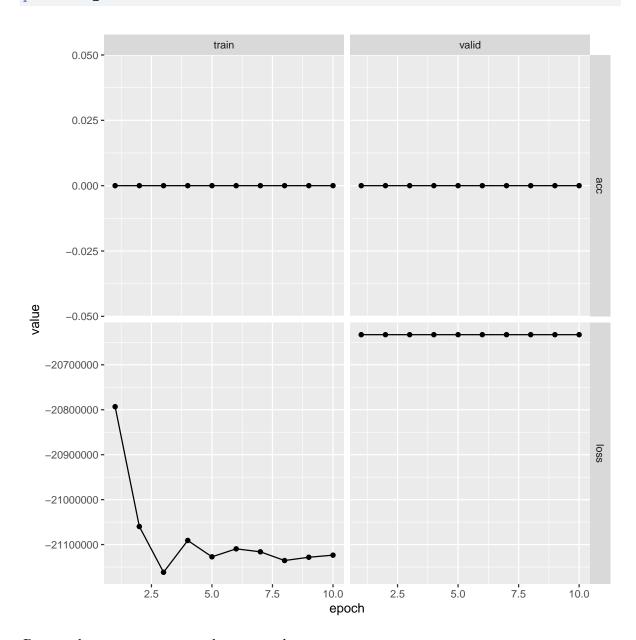
Fit a neural network model to predict the median\_house\_value using the same predictors as in 1.4. Use the model.matrix function to create the covariate matrix and luz package for fitting the network with 32, 16, 8 nodes in each of the three hidden layers.

```
covariate_matrix <- model.matrix(median_house_value ~ 0 + ., data = df_train)</pre>
nnet_fit <- NNet %>%
  setup(loss = nn_bce_loss(), opt = optim_adam, metrics = list(luz_metric_accuracy())) %>%
  set_hparams(p = 13, q1 = 32, q2 = 16, q3 = 8) %>%
  set_opt_hparams(lr = 0.02) %>%
  fit(data = list(covariate_matrix, df_train %>% select(median_house_value) %>% as.matrix),
      valid_data = list(model.matrix(median_house_value ~ 0 + ., data = df_test), df_test %
      epochs = 10,
      verbose = TRUE)
Epoch 1/10
Train metrics: Loss: -20793198 - Acc: 0
Valid metrics: Loss: -20633008 - Acc: 0
Epoch 2/10
Train metrics: Loss: -21059852 - Acc: 0
Valid metrics: Loss: -20633008 - Acc: 0
Epoch 3/10
Train metrics: Loss: -21161658 - Acc: 0
Valid metrics: Loss: -20633008 - Acc: 0
Epoch 4/10
Train metrics: Loss: -21090882 - Acc: 0
Valid metrics: Loss: -20633008 - Acc: 0
Epoch 5/10
Train metrics: Loss: -21127252 - Acc: 0
Valid metrics: Loss: -20633008 - Acc: 0
Epoch 6/10
Train metrics: Loss: -21109514 - Acc: 0
Valid metrics: Loss: -20633008 - Acc: 0
Epoch 7/10
Train metrics: Loss: -21116116 - Acc: 0
Valid metrics: Loss: -20633008 - Acc: 0
Epoch 8/10
Train metrics: Loss: -21135516 - Acc: 0
Valid metrics: Loss: -20633008 - Acc: 0
Epoch 9/10
Train metrics: Loss: -21128248 - Acc: 0
Valid metrics: Loss: -20633008 - Acc: 0
Epoch 10/10
Train metrics: Loss: -21123570 - Acc: 0
```

Valid metrics: Loss: -20633008 - Acc: 0

Plot the results of the training and validation loss and accuracy.

# plot(nnet\_fit)



Report the root mean squared error on the test set.

nnet\_predictions <- predict(nnet\_fit, model.matrix(median\_house\_value ~ 0 + ., data = df\_tes</pre>

```
nnet_rmse <- rmse(nnet_predictions, df_test$median_house_value)</pre>
nnet_rmse
```

#### [1] 236302.8



#### 🔔 Warning

Remember to use the as\_array() function to convert the predictions to a vector of numbers before computing the RMSE with rmse()

#### 1.9 (5 points)

Summarize your results in a table comparing the RMSE for the different models. Which model performed best? Why do you think that is?

```
rmse_df <- data.frame(Model = c("Linear Regression", "Decision Tree", "SVM", "Neural Network</pre>
rmse_df
```

```
Model
                         RMSE
1 Linear Regression 69082.77
2
     Decision Tree
                     74346.43
3
                SVM 61064.47
4
     Neural Network 236302.83
```

The SVM performed the best. This is likely because this is not a big data set, and we have not fine tuned the hyper parameters for the neural network.

#### Question 2



50 points

Spam email classification

The data folder contains the spam.csv dataset. This dataset contains features extracted from a collection of spam and non-spam emails. The objective is to classify the emails as spam or non-spam.

2.1 (2.5 points)

Read the data file as a tibble in R. Preprocess the data such that:

- 1. the variables are of the right data type, e.g., categorical variables are encoded as factors
- 2. all column names to lower case for consistency
- 3. Any observations with missing values are dropped

```
library(readr)
df2 <- read_csv("data/spambase.csv")

Rows: 4601 Columns: 58
-- Column specification ------
Delimiter: ","
db1 (58): word_freq_1, word_freq_2, word_freq_3, word_freq_4, word_freq_5, w...

i Use `spec()` to retrieve the full column specification for this data.
i Specify the column types or set `show_col_types = FALSE` to quiet this message.

df2 <- df2 %>%
  mutate(across(where(is.character), as.factor)) %>%
  rename_all(tolower) %>%
  na.omit()
```

2.2 (2.5 points)

Split the data df into df\_train and df\_split using test\_ind in the code below:

```
set.seed(42)
test_ind <- sample(
   1:nrow(df),
   floor( nrow(df)/10 ),
   replace=FALSE
)

df_train2 <- df2[-test_ind,]
df_test2 <- df2[test_ind,]
df_train2 <- na.omit(df_train2)
df_test2 <- na.omit(df_test2)</pre>
```

Complete the overview function which returns a data frame with the following columns: accuracy, error, false positive rate, true positive rate, between the true true\_class and the predicted pred\_class for any classification model.

```
overview <- function(predicted, expected) {</pre>
      confusion_matrix <- table(predicted,expected)</pre>
    accuracy <- (confusion_matrix[1] + confusion_matrix[4]) / length(expected) * 100 #
    error <- 100 - accuracy # Insert your code here
    total_false_positives <- confusion_matrix[3]</pre>
    total_true_positives <- confusion_matrix[4]</pre>
    total_false_negatives <- confusion_matrix[2]
    total_true_negatives <- confusion_matrix[1]</pre>
    false positive rate <- total false positives / (total false positives + total true negat
    false_negative_rate <- total_false_negatives / (total_false_negatives + total_true_posit
    return(
    data.frame(
    accuracy = accuracy,
    error=error,
    false_positive_rate = false_positive_rate,
    false_negative_rate = false_negative_rate
  )
}
```

#### 2.3 (5 points)

Fit a logistic regression model to predict the spam variable using the remaining predictors. Report the prediction accuracy on the test set.

```
glm_fit <- glm(spam ~., data = df_train2)
spam_predictions <- predict(glm_fit, newdata = df_test2, type = "response")
predicted_factors <- ifelse(spam_predictions > 0.5, 1, 0)
glm_classes <- overview(predicted_factors, df_test2$spam)
glm_classes</pre>
```

```
accuracy error false_positive_rate false_negative_rate 1 89.1258 10.8742 0.1360759 0.05228758
```

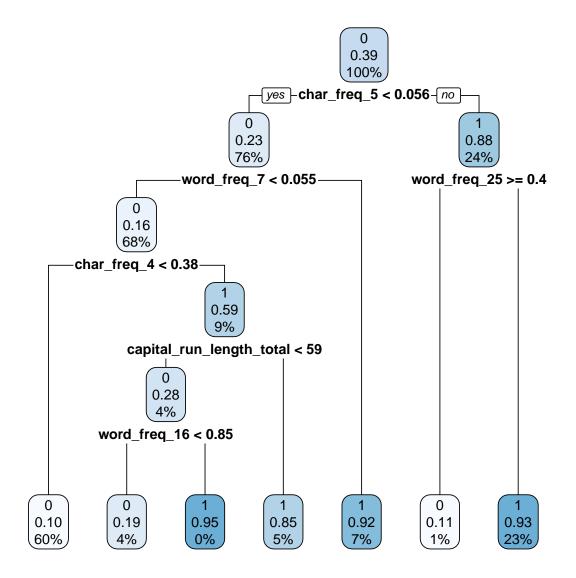
2.4 (5 points)

Fit a decision tree model to predict the spam variable using the remaining predictors. Use the rpart() function and set the method argument to "class".

```
rpart_classes <- rpart(spam ~. , data = df_train2, method = "class")
rpart_predictions2 <- predict(rpart_classes, newdata = df_test2, type = "class")</pre>
```

Visualize the decision tree using the rpart.plot() function.

```
rpart.plot(rpart_classes, box.palette = "Blues")
```



Report the prediction accuracy on the test set.

```
rpart_classes <- overview(rpart_predictions2, df_test2$spam)
rpart_classes</pre>
```

```
accuracy error false_positive_rate false_negative_rate 1 90.40512 9.594883 0.1066667 0.07692308
```

### 2.5 (5 points)

Fit a support vector machine model to predict the spam variable using the remaining predictors. Use the sym() function and use any kernel of your choice. Remember to set the type argument to "C-classification" if you haven't already converted spam to be of type factor.

```
svm_fit2 <- svm(spam ~. , data = df_train2, type = "C-classification")</pre>
```

Report the prediction accuracy on the test set.

```
svm_predictions2 <- predict(svm_fit2, newdata = df_test2, type = "class")</pre>
svm_classes <- overview(svm_predictions2, df_test2$spam)</pre>
svm_classes
```

```
accuracy
              error false_positive_rate false_negative_rate
1 92.96375 7.036247
                             0.07241379
                                                  0.06703911
```

### 2.6 (25 points)

Using the same neural network architecture as in 1.9, fit a neural network model to predict the spam variable using the remaining predictors.

#### A Classification vs. Regression

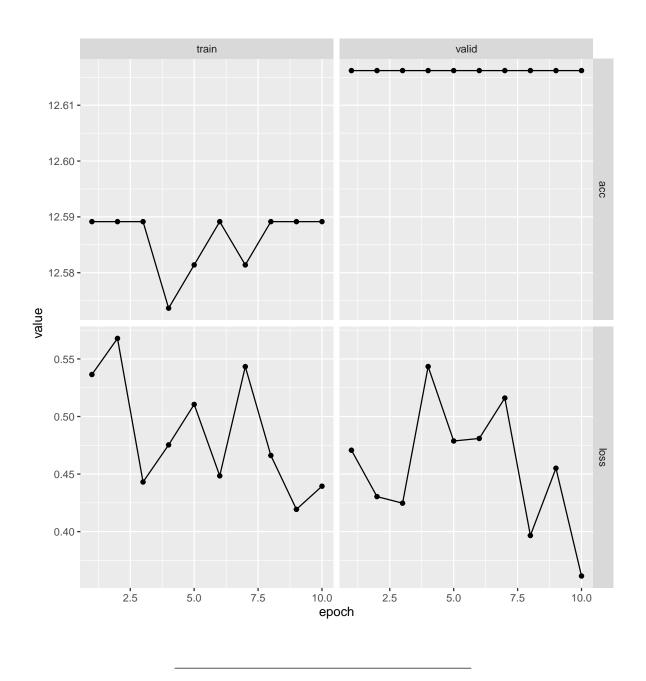
Note that the neural network in Q 1.9 was a regression model. You will need to modify the neural network architecture to be a classification model by changing the output layer to have a single node with a sigmoid activation function.

Use the model.matrix function to create the covariate matrix and luz package for fitting the network with 32, 16, 8 nodes in each of the three hidden layers.

```
covariate_matrix2 <- model.matrix(spam ~ 0 + ., data = df_train2)</pre>
nnet fit2 <- NNet %>%
  setup(loss = nn_bce_loss(), opt = optim_adam , metrics = list(luz_metric_accuracy())) %>%
  set_hparams(p = ncol(covariate_matrix2), q1 = 32, q2 = 16, q3 = 8) %>%
  set_opt_hparams(lr = 0.02) %>%
  fit(data = list(covariate_matrix2, df_train2 %>% select(spam) %>% as.matrix),
      valid_data = list(model.matrix(spam ~ 0 + ., data = df_test2), df_test2 %>% select(span)
      epochs = 10, verbose = TRUE)
```

```
Epoch 1/10
Train metrics: Loss: 0.5365 - Acc: 12.5891
Valid metrics: Loss: 0.4707 - Acc: 12.6162
Epoch 2/10
Train metrics: Loss: 0.5678 - Acc: 12.5891
Valid metrics: Loss: 0.4304 - Acc: 12.6162
Epoch 3/10
Train metrics: Loss: 0.4431 - Acc: 12.5891
Valid metrics: Loss: 0.4247 - Acc: 12.6162
Epoch 4/10
Train metrics: Loss: 0.4753 - Acc: 12.5736
Valid metrics: Loss: 0.5434 - Acc: 12.6162
Epoch 5/10
Train metrics: Loss: 0.5105 - Acc: 12.5814
Valid metrics: Loss: 0.4787 - Acc: 12.6162
Epoch 6/10
Train metrics: Loss: 0.4485 - Acc: 12.5891
Valid metrics: Loss: 0.4809 - Acc: 12.6162
Epoch 7/10
Train metrics: Loss: 0.5433 - Acc: 12.5814
Valid metrics: Loss: 0.5161 - Acc: 12.6162
Epoch 8/10
Train metrics: Loss: 0.4661 - Acc: 12.5891
Valid metrics: Loss: 0.3966 - Acc: 12.6162
Epoch 9/10
Train metrics: Loss: 0.4193 - Acc: 12.5891
Valid metrics: Loss: 0.4551 - Acc: 12.6162
Epoch 10/10
Train metrics: Loss: 0.4395 - Acc: 12.5891
Valid metrics: Loss: 0.3614 - Acc: 12.6162
```

### plot(nnet\_fit2)



# 2.7 (5 points)

Summarize your results in a table comparing the accuracy metrics for the different models.

nnet\_predictions <- predict(nnet\_fit2, model.matrix(spam ~ 0 + ., data = df\_test2)) %>% as.a.
nn\_classes2 <- overview(nnet\_predictions, df\_test2\$spam)</pre>

```
summary <- data.frame(Model = c("Logistic regression", "Decision tree", "SVM", "Neural Netwo
summary
```

```
Model
                        Accurage
1 Logistic regression 89.1257996
2
        Decision tree 90.4051173
3
                  SVM 92.9637527
4
       Neural Network 0.4264392
```

If you were to choose a model to classify spam emails, which model would you choose? Think about the context of the problem and the cost of false positives and false negatives.

The neural network had the lowest loss value, thus I would think to pick the neural network; however, the sym model had a much higher accuracy on the test data compared to the nnet. Thus, the svm would be more likely to be used if the test data is not biased.

### Question 3



Three spirals classification

To better illustrate the power of depth in neural networks, we will use a toy dataset called the "Three Spirals" data. This dataset consists of two intertwined spirals, making it challenging for shallow models to classify the data accurately.



⚠ This is a multi-class classification problem

The dataset can be generated using the provided R code below:

```
generate_three_spirals <- function(){</pre>
  set.seed(42)
  n < -500
  noise \leftarrow 0.2
  t \leftarrow (1:n) / n * 2 * pi
  x1 <- c(
      t * (sin(t) + rnorm(n, 0, noise)),
      t * (sin(t + 2 * pi/3) + rnorm(n, 0, noise)),
```

```
t * (sin(t + 4 * pi/3) + rnorm(n, 0, noise))
)

x2 <- c(
    t * (cos(t) + rnorm(n, 0, noise)),
    t * (cos(t + 2 * pi/3) + rnorm(n, 0, noise)),
    t * (cos(t + 4 * pi/3) + rnorm(n, 0, noise))
)

y <- as.factor(
    c(
        rep(0, n),
        rep(1, n),
        rep(2, n)
)

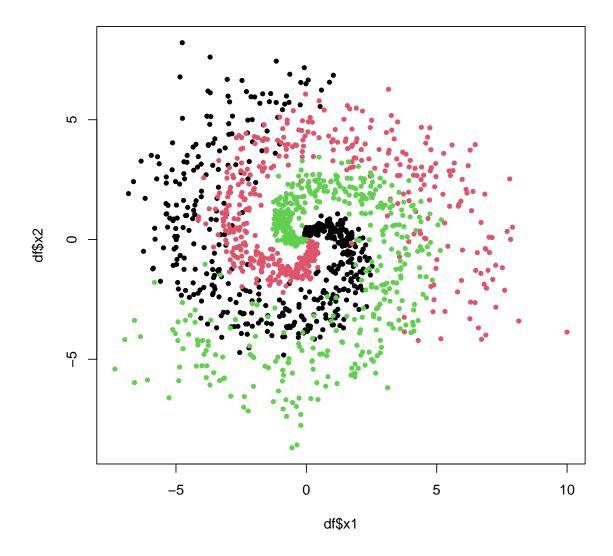
return(tibble::tibble(x1=x1, x2=x2, y=y))
}</pre>
```

## 3.1 (5 points)

Generate the three spirals dataset using the code above. Plot  $x_1$  vs  $x_2$  and use the y variable to color the points.

```
df <- generate_three_spirals()

plot(
   df$x1, df$x2,
   col = df$y,
   pch = 20
)</pre>
```



Define a grid of 100 points from -10 to 10 in both  $x_1$  and  $x_2$  using the expand.grid(). Save it as a tibble called df\_test.

```
x_values <- seq(from = -10, to = 10, length.out = 100)
grid <- expand.grid(x1 = x_values, x2 = x_values) # Insert your code here

n <- nrow(grid)
y <- as.factor(c(rep(0, n), rep(1, n), rep(2, n)))</pre>
```

```
# Combine x1, x2, and y into a tibble
df_test2 <- cbind(grid, y)
df_test2 <- as_tibble(df_test2)</pre>
```

## 3.2 (10 points)

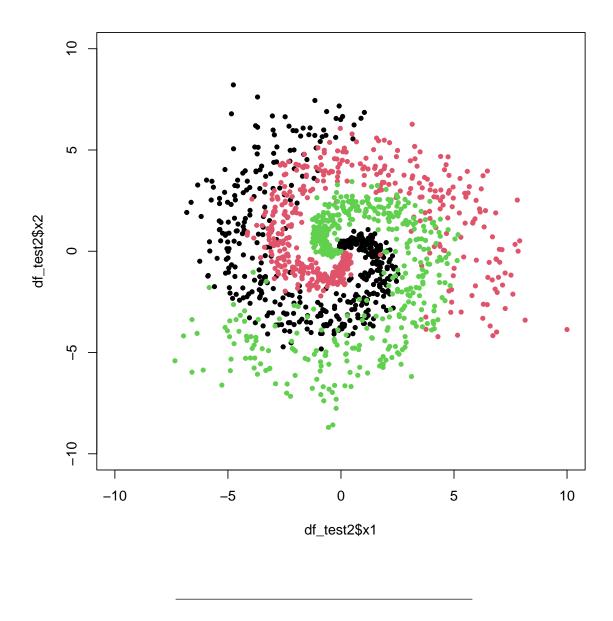
Fit a classification tree model to predict the y variable using the x1 and x2 predictors, and plot the decision boundary.

```
rpart_fit <- rpart(y ~ x1 + x2, data = df, method = "class") # Insert your code here
rpart_classes <- predict(rpart_fit, df_test2)</pre>
```

Plot the decision boundary using the following function:

```
plot_decision_boundary <- function(predictions){
  plot(
    df_test2$x1, df_test2$x2,
    col = predictions,
    pch = 0
)
  points(
    df$x1, df$x2,
    col = df$y,
    pch = 20
)
}</pre>
```

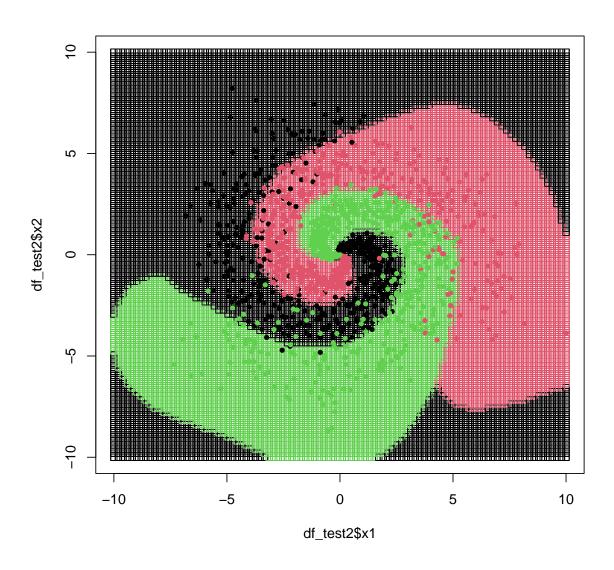
```
plot_decision_boundary(rpart_classes)
```



# 3.3 (10 points)

Fit a support vector machine model to predict the y variable using the x1 and x2 predictors. Use the svm() function and use any kernel of your choice. Remember to set the type argument to "C-classification" if you haven't converted y to be of type factor.

```
svm_fit <- svm(y ~ x1 + x2, data = df, type = "C-classification")
svm_classes <- predict(svm_fit, newdata = df_test2)
plot_decision_boundary(svm_classes)</pre>
```



## ⚠ Instructions

For the next questions, you will need to fit a series of neural networks. In all cases, you can:

- set the number of units in each hidden layer to 10
- set the output dimension o to 3 (remember this is multinomial classification)
- use the appropriate loss function for the problem (**not nn\_bce\_loss**)
- set the number of epochs to 50
- fit the model using the luz package

You can use any optimizer of your choice, but you will need to tune the learning rate for each problem.

### 3.4 (10 points)

Fit a neural network with 1 hidden layer to predict the y variable using the x1 and x2 predictors.

```
NN1 <- nn_module(</pre>
  initialize = function(p, q1, o){
    self$hidden1 <- nn_linear(p,q1)</pre>
    self$output <- nn_linear(q1, o)</pre>
    self$activation <- nn_relu()</pre>
  },
  forward = function(x){
    x %>%
      self$hidden1() %>%
      self$activation() %>%
      self$output()
  }
)
df_test2 <- df_test2 %>% mutate(y = as.numeric(y))
fit_1 <- NN1 %>%
    loss = nn_cross_entropy_loss(), opt = optim_adam, metric = luz_metric_accuracy()
  ) %>%
  set_hparams(
    p = ncol(df), q1 = 10, o = 3
```

In order to generate the class predictions, you will need to use the predict() function as follows

```
test_matrix <- df_test2 %>% select(x1, x2) %>% as.matrix
fit_1_predictions <- predict(fit_1, test_matrix) %>%
    as.array()

plot_decision_boundary(overview(fit_1_predicitons, df_test2$y))
```

Plot the results using the plot\_decision\_boundary() function.

\_\_\_\_

### 3.5 (10 points)

Fit a neural network with 0 hidden layers to predict the y variable using the x1 and x2 predictors.

```
NNO <- nn_module(
  initialize = function(p, o){
    self$output <- nn_linear(p, o)
    self$activation <- nn_relu()
    self$sigmoid <- nn_sigmoid()
},
forward = function(x){
    x %>%
```

```
self$hidden1() %>%
      self$activation() %>%
      self$output() %>%
      self$sigmoid()
  }
)
fit_0 <- NNO %>%
  setup(loss = nn_cross_entropy_loss(), opt = optim_adam, metric = luz_metric_accuracy()) %>
  set_hparams( p = ncol(df), o = 3) %>%
  set_opt_hparams(lr = 0.02) %>%
  fit(data = list(
     df %>% select(x1,x2) %>% as.matrix, df %>% select(y)
    ),
    valid = list(
            df_test2 %>% select(x1,x2) %>% as.matrix, df_test2 %>% select(y)
    epochs = 50,
    verbose = TRUE)
```

Plot the results using the plot\_decision\_boundary() function.

```
test_matrix <- df_test2 %>% select(x1, x2) %>% as.matrix

fit_0_predictions <- predict(fit_0, test_matrix) %>%
    as.array()

plot_decision_boundary(overview(fit_0_predicitons, df_test2$y))
```

#### 3.6 (10 points)

Fit a neural network with 2 hidden layers to predict the y variable using the x1 and x2 predictors.

```
NN2 <- nn_module(
  initialize = function(p, q1, q2, o){
    self$hidden1 <- nn_linear(p,q1)
    self$hidden2 <- nn_linear(q1, q2)
    self$output <- nn_linear(q2,o)</pre>
```

```
self$activation <- nn_relu()</pre>
    self$sigmoid <- nn_sigmoid()</pre>
  },
  forward = function(x){
    x %>%
      self$hidden1() %>%
      self$activation() %>%
      self$hidden2() %>%
      self$activation() %>%
      self$output() %>%
      self$sigmoid()
  }
fit_2 <- NN3 %>%
  setup(loss = nn_cross_entropy_loss(), opt = optim_adam, metric = luz_metric_accuracy()) %>'
  set_hparams(p = ncol(df), q1 = 10, q2 = 10, o = 3) %>%
  set_opt_params(lr = 0.02) %>%
  fit(data = list(
      model.matrix(y ~ 0 + ., data = df), df %>% select(y) %>% as.matrix
    ),
    valid = list(
            model.matrix(y ~ 0 + ., data = df_test2), df_test2 %>% select(y) %>% as.matrix
    epochs = 50,
    verbose = TRUE)
```

Plot the results using the plot\_decision\_boundary() function.

```
fit_2_predictions <- predict(fit_2, test_matrix) %>%
    as.array()

plot_decision_boundary(overview(fit_2_predicitons, df_test2$y))
```

#### 3.7 (5 points)

What are the differences between the models? How do the decision boundaries change as the number of hidden layers increases?

## i Session Information

Print your R session information using the following command

#### sessionInfo()

```
R version 4.3.1 (2023-06-16 ucrt)
```

Platform: x86\_64-w64-mingw32/x64 (64-bit)
Running under: Windows 11 x64 (build 22631)

Matrix products: default

#### locale:

- [1] LC\_COLLATE=English\_United States.utf8
- [2] LC\_CTYPE=English\_United States.utf8
- [3] LC\_MONETARY=English\_United States.utf8
- [4] LC\_NUMERIC=C
- [5] LC\_TIME=English\_United States.utf8

time zone: America/New\_York
tzcode source: internal

## attached base packages:

[1] stats graphics grDevices datasets utils methods base

## other attached packages:

[1]	luz_0.4.0	torch_0.12.0	e1071_1.7-14	<pre>rpart.plot_3.1.2</pre>
[5]	rpart_4.1.23	caret_6.0-94	lattice_0.21-8	ggplot2_3.5.0
[9]	corrplot_0.92	magrittr_2.0.3	broom_1.0.5	purrr_1.0.2
Г137	tidvr 1.3.1	readr 2.1.5	dplvr 1.1.4	tibble 3.2.1

#### loaded via a namespace (and not attached):

[1]	tidyselect_1.2.1	timeDate_4032.109	farver_2.1.1
[4]	fastmap_1.1.1	pROC_1.18.5	digest_0.6.35
[7]	timechange_0.3.0	lifecycle_1.0.4	ellipsis_0.3.2
[10]	processx_3.8.4	survival_3.5-5	compiler_4.3.1
[13]	progress_1.2.3	rlang_1.1.3	tools_4.3.1
[16]	utf8_1.2.4	yaml_2.3.8	data.table_1.15.2
[19]	knitr_1.45	labeling_0.4.3	<pre>prettyunits_1.2.0</pre>

[22] bit_4.0.5	plyr_1.8.9	withr_3.0.0
[25] nnet_7.3-19	grid_4.3.1	stats4_4.3.1
[28] fansi_1.0.6	colorspace_2.1-0	future_1.33.2
[31] globals_0.16.3	scales_1.3.0	iterators_1.0.14
[34] MASS_7.3-60	zeallot_0.1.0	cli_3.6.2
[37] crayon_1.5.2	rmarkdown_2.26	generics_0.1.3
[40] future.apply_1.11.2	reshape2_1.4.4	tzdb_0.4.0
[43] proxy_0.4-27	stringr_1.5.1	splines_4.3.1
[46] parallel_4.3.1	coro_1.0.4	vctrs_0.6.5
[49] hardhat_1.3.1	Matrix_1.5-4.1	jsonlite_1.8.8
[52] callr_3.7.6	hms_1.1.3	bit64_4.0.5
[55] listenv_0.9.1	foreach_1.5.2	gower_1.0.1
[58] recipes_1.0.10	glue_1.7.0	parallelly_1.37.1
[61] ps_1.7.6	codetools_0.2-19	lubridate_1.9.3
[64] stringi_1.8.3	gtable_0.3.4	munsell_0.5.0
[67] pillar_1.9.0	htmltools_0.5.8	ipred_0.9-14
[70] lava_1.8.0	R6_2.5.1	vroom_1.6.5
[73] evaluate_0.23	backports_1.4.1	renv_1.0.3
[76] class_7.3-22	Rcpp_1.0.12	nlme_3.1-162
[79] prodlim_2023.08.28	xfun_0.43	fs_1.6.3
[82] pkgconfig_2.0.3	ModelMetrics_1.2.2.2	!